

Kihwan Kim

Software Engineer · 7 years — deploying LLM solutions with the customers and teams who use them

juljin1875@gmail.com
kihwan.kim ·
github.com/simulacre7
Seoul, South Korea

Software engineer with 7 years of experience turning ambiguous business problems into deployed AI solutions. At Levvels I embedded with an operations team, scoped how they carried state across several internal tools to make one decision, and shipped an agent-based automation that cut hours of manual investigation to minutes — then fed what I learned in the field back into the product. I build with AI coding tools daily (Claude Code, under a builder/reviewer workflow with output verification), and my background in HCI and Visual Analytics keeps me focused on making LLM output something people can trust and act on. I own problems end to end — from a customer conversation to a working interface — and I am at home in Python and the Unix command line.

Highlights

- **Agentic AI deployment, field to product** — defined the problem with an operations team, deployed an agent that turns natural-language requests into browser actions across internal tools, and channeled field insight back into product improvements (Levvels, PageAgent)
- **AI-assisted development** — ship daily with Claude Code under a builder/reviewer workflow: the builder writes, a separate model reviews, and behavior is pinned by tests before merge — the review caught a P0 just before submission in a live coding test
- **Making LLM output usable for non-experts** — built codeless studio flows and explainable-AI visualizations so enterprise users without an ML background could operate an AutoML platform and act on its predictions (TmaxEnterprise)
- **Owning across the stack** — contributed 10 merged PRs to an open-source game, from the TypeScript client into the Rust server; I can stand up a usable interface as well as work the backend

Experience

Levvels — Software Engineer

Jul 2025 – Present

AI-native product development and internal operations automation — scoping problems in the field and turning them into deployed solutions.

PageAgent-based Operations Automation agentic AI · deployment

- Embedded with the operations team to define the real problem: staff carried state across customer chat, delivery tracking, internal admin, and email to make a single decision
- Designed and deployed an agent-based automation that turned hours of manual investigation into minutes, connecting the needed context into one verifiable result view
- Extended Alibaba's open-source PageAgent for a Chrome Managed Profile and added a Slack human-in-the-loop hand-off the base tool lacked, so a stuck automation fails loudly to a person instead of returning a wrong result silently
- Fed field insights back to the product team, and separated LLM output from reproducible action sequences so operators can re-verify each conclusion against its source

AI Character Chat LLM streaming

- Separated in-progress streamed responses from persisted history; designed the send / complete / fail / regenerate state machine for interrupt-and-retry-heavy LLM responses
- Built authoring and review tooling for character content and reconciled specs across product docs, Figma, and Swagger

React · Next.js · TypeScript · Node.js · SSE · Playwright · Python · Claude Code · PageAgent · Generative UI

RIDI Corp. — Frontend Engineer

May 2022 – Jul 2025

Kept operational UI changes and experiments off the deployment critical path on a large, high-traffic content platform.

Server-driven UI & large-scale rendering

- Rendered server-declared UI on the client to decouple screen changes from deploys; migrated PHP pages toward React incrementally to lower risk
- Optimized large-list rendering with windowing and dynamic-height measurement; used Redis to cache server-driven UI payloads and reduce backend load; tracked real-user errors with Sentry

React · Next.js · TypeScript · Server-driven UI · Redis · Sentry · Virtualization

TmaxEnterprise — Research Engineer

Feb 2020 – May 2022

Made complex AI and analytics usable for non-experts — usability and explainability first.

AutoML Platform explainable AI · Python

- Built a codeless studio so users without model-development experience could build and operate models
- Developed explainable-AI visualizations that show the reasoning beside the prediction, so a non-expert has a basis to act

React · TypeScript · Material-UI · Python

Open Source

OpenMMO — Contributor

Jul 2026 – Present

- 10 PRs merged, owning changes end-to-end from the TypeScript client into the **Rust server** (cache eviction, account bans, trade validation) while learning a new backend language on a live codebase
- Paired with a coding agent under a builder/reviewer workflow, pinning behavior with tests before submission

vite-plus — Contributor

Aug 2026

- Fixed a shell-integration gap in Rust and covered the four shells' templates with snapshot tests (#2508, merged) — VoidZero, the team behind Vite and Vitest

Education

Ulsan National Institute of Science and Technology (UNIST)

2013 – 2020

M.S. in Computer Science (HCI · Visual Analytics; human-AI collaboration).

B.S. in Technology Management (CS multi-disciplinary major).

- “An Empirical Analysis on Transparent Algorithmic Exploration in Recommender Systems” — **CoRR 2021**
· arxiv.org/abs/2108.00151
- “Modeling Exploration/Exploitation Decisions through Mobile Sensing for Understanding Mechanisms of Addiction” — **MobiSys '19** · dl.acm.org/doi/10.1145/3307334.3328599

Full CV, case studies, and portfolio at kihwan.kim.